

Round 2 Challenge: Artifact Hub

Context

Teams across the company are increasingly using AI tools — Claude, Gamma, GPT, Midjourney — to generate mockups, presentations, reports, and documentation. The output is useful, but the lifecycle after generation is messy: files end up in blob storage accessed via CLI commands, shared via expiring URLs pasted into Slack, with feedback scattered across threads. There's no way to browse what exists, no structured feedback mechanism, and no access control beyond URL expiry.

This is a real problem we face today. We want you to solve it.

The Challenge

Build **Artifact Hub** — a platform for publishing, browsing, reviewing, and sharing AI-generated content. It should feel like a product, not a homework assignment.

Core capabilities

- **Publish** artifacts (HTML, images, PDFs — at minimum) with metadata (title, description, tags/categories)
- **Browse** a gallery or catalog of published artifacts
- **Share** artifacts with configurable access (time-limited links, at minimum)
- **Feedback** — structured way to leave and view comments on artifacts
- **MCP server** — so that Claude Desktop, CoPilot, or other MCP clients can publish and manage artifacts conversationally

Making it intelligent

The MCP server already makes this an AI-native platform — an LLM can publish, search, and manage artifacts through natural conversation. Beyond that, find at least

one place where AI meaningfully improves the experience. Some possibilities (you're not limited to these):

- Summarizing feedback across multiple reviewers
- Auto-generating descriptions or tags from artifact content
- Natural language search across the catalog
- Smart review routing based on content analysis
- Something we haven't thought of

The best LLM features feel invisible — they make the product better without making "AI" the point. The worst feel bolted on. Use your judgment.

What we're evaluating

This is a Round 2 challenge. We already know you can research and write findings. Now we want to see you build and ship.

Dimension	What we're looking for
Product sense	Did you build something a team would actually use? Or a feature checklist? The choices you make about what to build, what to skip, and how it feels matter more than feature count.
User experience	Does the browsing, sharing, and feedback flow feel considered? Would a non-technical person understand how to use it?
Architecture	Is the system well-structured? Could someone else extend it? Are the boundaries clean?
MCP integration	Is the MCP server thoughtful — or just CRUD wrappers? Does it make the conversational workflow feel natural?
LLM integration	Is AI used where it genuinely helps? Does it feel like a product feature or a demo?
Deployment	Is the system hosted and accessible? Can a reviewer use it without setting up a local environment?

Dimension	What we're looking for
Engineering quality	TypeScript, clean code, reasonable error handling, evidence that the code works (tests, or at minimum, a system that demonstrably runs)

Deliverables

1. **Running system** at a publicly accessible URL
 2. **MCP server configuration** that a reviewer can add to Claude Desktop to interact with the platform
 3. **WRITEUP.md** in the repo root:
 - What you built and why (product decisions)
 - What you chose not to build and why
 - Architecture overview
 - How the MCP integration works
 - Where and why you used LLM capabilities
 - Deployment approach
 - What you'd do next with another week
 4. **Brief walkthrough** — either a screen recording (5 min max) or a written step-by-step showing the key flows
 5. **Claude Code session logs** (see below)
-

Sharing your Claude Code sessions

We evaluate how you work with AI tools, not just the output. Please include your Claude Code session logs:

1. Locate your session files: `ls ~/.claude/projects/` — find the directory matching your project path

2. The session files are `.jsonl` files inside that directory
3. Copy the relevant session files into your submission repo under a `claude-sessions/` folder
4. If you used Claude Code commands, agents, or MCP servers, include your `.claude/` project directory as well

If you used other AI tools (Cursor, Copilot, ChatGPT), note which tools you used and for what in your `WRITEUP.md`. There are no wrong answers here — we want to see the workflow.

Scope and time

Time box: 2 days. This is a build challenge, not a research spike. Prioritize a working system over comprehensive features. A polished core that works is better than a sprawling system that doesn't.

You have full discretion over technology choices, hosting, and scope prioritization. The evaluation criteria above tell you what matters — how you allocate your time across those dimensions is part of what we're assessing.

Starting points

- [Model Context Protocol specification](#)
 - [MCP TypeScript SDK](#)
 - [Claude Desktop MCP setup](#)
 - For hosting: Netlify, Vercel, Railway, Fly.io, Render, Azure — your call
-

Submission

Submit via the same mechanism as Round 1. Commit history matters — we review how you work, not just what you ship.